



Centralizing Claude Code & GitHub Copilot governance

One control plane for the two agentic developer platforms
your engineers actually use — set policy once, prove it
continuously, never touch the inference path.

✱ Claude Code

👤 GitHub Copilot

■ The agentic dev estate has no control tower

🔗 THE PROBLEM

Two platforms now write a large share of your code: **Claude Code** and **GitHub Copilot**. Each ships its own admin surface — different consoles, different APIs, different gaps. Today, proving “which MCP servers are allowed, who can skip code review, and what we’re spending” means stitching together a dozen screens and trusting that nobody changed a setting last Tuesday.

TODAY — FRAGMENTED

Claude Code settings live on each **developer’s laptop**

Copilot policy split across **org + enterprise** consoles

MCP allow-lists set **per-tool, per-person**

Code-review rules in **each repo** separately

Spend in **two billing portals**, no per-user view

No record of **what was true** 30 days ago



WITH AIR-TRAFFIC — ONE TOWER

One **policy-as-code** file for both platforms

MCP, hooks & review gates **pushed + locked**

Per-user spend & usage, **side by side**

Every change in **one audit stream**

Continuous drift proof for auditors

Zero inference proxies on the request path

2→1

platforms governed from a single control plane

0

inference proxies on the critical path — pure control + config

A-tier

code-review & MCP gates that are server-side, non-bypassable

30-day

rolling policy-vs-actual drift history, audit-ready

■ What Air-Traffic centralizes

🔗 THE CONTROL CATALOG

Air-Traffic governs each platform through whichever mechanism actually works — and labels it honestly. **VendorNative** drives the vendor’s admin API. **EnvManaged** pushes managed configuration into the developer environment and reads it back. **ProxyEnforced** is an optional add-on (off by default). **MonitorOnly** verifies & alerts. The tier badge (**A** server-side · **B** MDM-locked · **C** seed-only) tells you how strong the enforcement really is.

 Claude Code ANTHROPIC · managed-settings.json + Admin/Analytics API		
MCP servers — allow / deny	EnvManaged B	<code>allowedMcpServers</code> / <code>deniedMcpServers</code> / <code>allowManagedMcpServersOnly</code> — locked on MDM-managed devices
Shared hooks	EnvManaged B	<code>allowManagedHooksOnly</code> , <code>allowedHttpHookUrls</code> — org-injected pre/post hooks
Permissions / tool rules	EnvManaged B	<code>allowManagedPermissionRulesOnly</code> — managed rules override user & project
Skill / agent execution	EnvManaged B	<code>disableSkillShellExecution</code> , <code>disableAgentView</code>
Model assignment	EnvManaged B	restrict callable models via managed settings (raw-API path → optional gateway)
Per-user usage & spend	VendorNative	Anthropic Usage/Cost API + Claude Code Analytics (commits, PRs, lines, sessions); hard cap → MonitorOnly
Drift read-back	VendorNative	<code>claude doctor</code> reports each effective setting + its source

 GitHub Copilot GITHUB · org / enterprise AI Controls + REST		
MCP server access	EnvManaged A	registry-only allowlist, enforced server-side — user cannot add unlisted servers
Feature policies	EnvManaged A	Chat / CLI / Cloud Agent / Extensions on-off via Enterprise AI Controls
Content exclusions	VendorNative	REST <code>/orgs/{org}/copilot/content_exclusion</code> — repos & path globs
Code-review gate	EnvManaged A	branch-protection / rulesets: required human reviewers, enforced at merge
Test-coverage gate	EnvManaged A	required status checks (e.g. <code>coverage-90</code>) before an agent commit lands
Seat & AI-credit spend	VendorNative	seats + <code>billing/usage</code> API; 1 credit = \$0.01; hard token cap → MonitorOnly
Usage metrics	VendorNative	<code>/copilot/metrics</code> — 28-day rolling, by team / language / editor
Agentic audit	VendorNative	<code>agent_session.task</code> events, <code>actor:Copilot</code> (GA Feb 2026), commit→session, 180-day

■ Enforcement confidence, stated honestly

🔗 HOW STRONG IS “ENFORCED”?

“Governed” is not one thing. Air-Traffic grades every control by how hard it is to bypass — and never paints a soft control green.

A Server-side enforced — non-bypassable

GitHub / GitLab branch protection & rulesets (code-review + test-coverage gates), Copilot MCP registry-only allowlist, feature policies. Enforced on the vendor’s servers; a developer simply cannot override it locally. **Start here.**

B MDM-locked client config — enforced on managed devices

Claude Code `managed-settings.json` (MCP, hooks, permissions, skill exec) takes top precedence over user & project settings — but only when the file is present. Real enforcement requires the device to be MDM-enrolled (Jamf / Intune / Kandji).

C Seed + drift-detect only — monitored, not enforced


Local client settings on unmanaged devices. Air-Traffic distributes the config and **alerts on divergence**, but cannot prevent an override. The UI labels these plainly — they are never shown as “enforced.”

The one honest caveat. Tier-B controls (Claude Code managed settings) are genuine, top-precedence enforcement — conditional on MDM enrollment. Air-Traffic verifies MDM coverage before claiming Tier-B, and degrades to “monitored” with a drift alert on any unmanaged device. No overstatement, ever.

■ How it works — no proxy required

🔗 ARCHITECTURE

Air-Traffic is a **control plane + config distributor**, not a man-in-the-middle on your model calls. It drives each platform’s admin surface, pushes managed config into the environment, and reads the effective state back for drift — then normalizes everything into one observability stream.



What is not here: there is no gateway intercepting `/v1/messages` or Copilot completions. Spend is monitored from the vendors' usage APIs and capped softly; only if you ever need a hard, mid-request, cross-vendor cutoff do you switch on the optional gateway module — and nothing else depends on it.

■ Proof it's real

🔗 THE ARTIFACTS AIR-TRAFFIC EMITS

One policy declaration compiles into the exact managed artifacts each platform understands — pushed, then verified.

Claude Code — managed-settings.json (Tier-B, MDM-locked)

```

{
  "permissions": { "allowManagedPermissionRulesOnly": true },
  "allowManagedMcpServersOnly": true,
  "allowedMcpServers": ["filesystem", "git", "internal-db-readonly"],
  "deniedMcpServers": ["*"],
  "allowManagedHooksOnly": true,
  "hooks": { "PreToolUse": ["secret-scan"], "PostCommit": ["run-tests"] },
  "disableSkillShellExecution": true
}
// deployed via MDM → /Library/Application Support/ClaudeCode/managed-settings.json
// verified with: claude doctor (effective setting + source)

```

GitHub — branch protection / ruleset (Tier-A, server-side)

```
{
  "required_pull_request_reviews": { "required_approving_review_count": 1,
                                     "require_code_owner_reviews": true },
  "required_status_checks": { "strict": true,
                              "checks": [{ "context": "tests" },
                                           { "context": "coverage-90" }] },
  "enforce_admins": true
}
// PUT /repos/{org}/{repo}/branches/{branch}/protection
// → an agent commit cannot merge without human review + passing coverage
```

The control room

EXECUTIVE FLIGHT DECK — SCOPED TO THE TWO PLATFORMS

Air-Traffic · Flight Deck policy: Fintech · synced 12s ago				
PLATFORM	POLICY SYNC	TOP ENFORCEMENT	SPEND (MTD)	DRIFT
✱ Claude Code	✓ synced	MCP + hooks EnvManaged B	\$6,240	none
● GitHub Copilot	✓ synced	Review + MCP EnvManaged A	\$4,815	none
Platform Eng — laptops MDM-enrolled	42 / 44 managed monitored	2 unmanaged →		
Agentic actions audited (30d)	18,402 events · 0 unattributed			

Drilling into any row opens the change history: who set it, when, the before/after, and whether it's holding. Every cell traces back to a normalized audit event — the same ops-observation-batch/v1 stream a SIEM can ingest.

Why it matters

THE BOARD-LEVEL OUTCOME

Governance that proves itself

- **One policy-as-code source of truth** spanning Claude Code and Copilot — reviewed like any other code.
- **Continuous drift proof.** Every audit opens with 30 days of policy-vs-actual history, not a screenshot.
- **Shadow-MCP discovery.** Any unlisted MCP server or unmanaged device surfaces as drift, not a surprise.
- **Honest posture.** Tier badges mean leadership knows exactly what is enforced vs. monitored.

Spend & velocity, in one view

- **Per-user, cross-platform spend** — Claude Code tokens and Copilot AI-credits side by side.
- **Agent productivity** — commits, PRs, accepted suggestions, by team, normalized across both.
- **Soft caps now, hard caps optional** — monitor and alert today; switch on the gateway only if you need a mid-request cutoff.

1

policy file governs
both platforms

100%

of agentic actions
audited where the
platform exposes
them

A+B

tiers give real
enforcement — server-
side & MDM-locked

5 min

to apply an industry
baseline to the whole
estate

Air-Traffic · Enterprise AI Control Plane & Observability — Claude Code x GitHub Copilot · June 2026
Vendor marks: official Claude Code & GitHub Copilot glyphs (simple-icons). Air-Traffic brand: option-4 logo pack.